

Chapter 1

Instruction Set Manual

1.1 Compute Processor Instruction Set

Refer to Section 1.2 for concrete definitions of semantic helper functions such as *sign-extend-16-to-32* and *rotate-left*.

1.1.1 Register Conventions

The compute processor uses register conventions similar to those used in MIPS microprocessors. Procedures cannot rely on *caller-saved* registers retaining their values upon a procedure call. Procedures must restore the initial values of *callee-saved* registers before returning to their caller. The conventions are shown below.

| Register Number | Assembly Alias | Saved by | Description |
|-----------------|----------------|----------|---|
| \$0 | \$at | n/a | Always has value zero. |
| \$1 | | caller | Assembler temporary clobbered by some assembler operations. |
| \$2..\$3 | | caller | First and second words of return value, respectively. |
| \$4..\$7 | | caller | First 4 arguments of function. |
| \$8..\$15 | | caller | General registers. |
| \$16..\$23 | | callee | General registers. |
| \$24 | \$cst[i/o] | n/a | Static Network input/output port. |
| \$25 | \$cgn[i/o] | n/a | General Dynamic Network input/output port. |
| \$26 | \$csti2 | n/a | Static Network input port #2. |
| \$27 | \$cmn[i/o] | n/a | Memory Dynamic Network input/output port. |
| \$28 | \$gp | callee | Global pointer. Points to start of tile's code and static data. |
| \$29 | \$sp | callee | Stack pointer. Stack grows towards lower addresses. |
| \$30 | | callee | General register. |
| \$31 | | caller | Link register. Saves return address for function call. |

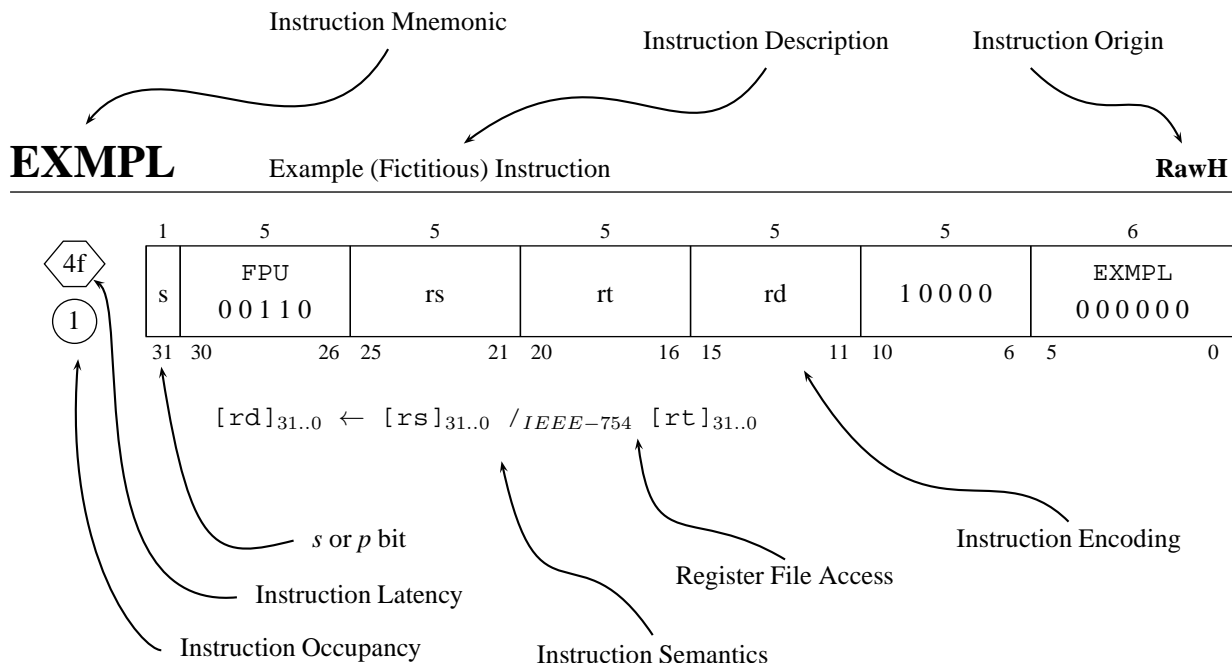
⁰Based on a Template provided by Prof. Michael Taylor of UC San Diego. Free for general use as long as this notice remains here.

1.1.2 Compute Processor Instruction Template

Shown below is an example instruction listing. The *instruction occupancy* is the number of compute processor issue cycles that are occupied by the instruction. Subsequent instructions must wait for this number of cycles before issuing. The *instruction latency* is the total number of cycles that must pass before a subsequent dependent instruction can issue. Suffixes of *d* and *f* indicate that an instruction uses the integer and floating-point divide units, respectively, for that number of cycles. Subsequent instructions that require a particular unit will stall until that unit is free. A suffix of *b* means that the instruction has an additional 3 cycles of occupancy on a branch misprediction. An occupancy of *c* means that the instruction takes at least 13 cycles if a cache line is evicted, 5 cycles if only an invalidation occurs, otherwise 1 cycle.

The *s* or *p* bits in the instruction encoding specify respectively whether 1) an instruction's output will be copied to *cst0* in addition to the destination register, or 2) whether a branch is predicted taken or not.


Generally, the Raw compute processor attempts to inherit the MIPS instruction set mnemonics to the extent that it reduces the learning curve for new users of the system. However, the underlying instruction semantics have been "cleaned up"; for instance, interlocks have been added for load, branch, multiply and divide instructions (reducing the need to insert *nops*), and the FPU uses the same register set as the ALU. To this end, the *instruction origin* specifies whether the instruction semantics are very similar the MIPS instruction of the same name ("MIPS"), whether they are specific to the Raw architecture ("'" or Raw), or to the Raw architecture extended with hardware instruction caching ("RawH"). Of course, the instruction encodings (including the presence of *s* and *p* bits) are completely different from MIPS.



ADDIU

Add Immediate


MIPS


$$\text{sim16}_{31..0} \leftarrow (\text{sign-extend-16-to-32 } \text{sim16})$$
$$[\text{rt}]_{31..0} \leftarrow \{ [\text{rs}]_{31..0} + \text{sim16} \}_{31..0}$$

ADDU

Add


MIPS


$$[\text{rd}]_{31..0} \leftarrow \{ [\text{rs}]_{31..0} + [\text{rt}]_{31..0} \}_{31..0}$$

AND

And Bitwise


MIPS


$$[\text{rd}]_{31..0} \leftarrow [\text{rs}]_{31..0} \& [\text{rt}]_{31..0}$$

ANDI


And Bitwise Immediate

MIPS


$$[\text{rt}]_{31..16} \leftarrow [\text{rs}]_{31..16}$$
$$[\text{rt}]_{15..0} \leftarrow [\text{rs}]_{15..0} \& \text{uimm16}$$

AUI

Add Upper Immediate


$$[rt]_{31..16} \leftarrow \{ [rs]_{31..16} + uimm16 \}_{15..0}$$
$$[rt]_{15..00} \leftarrow [rs]_{15..00}$$

B


Branch Unconditional (Assembly Macro)

① $b <label>$

$$PC_{31..00} \leftarrow <label>$$


BEQ

Branch if equal


$$\text{if } ([rs] == [rt])$$
$$PC_{31..02} \leftarrow \{ PC_{31..02} + (sign-extend-16-to-30 \text{ boffs16}) \}_{29..00}$$
$$PC_{01..00} \leftarrow 0$$


BGEZ

Branch if greater than or equal to zero (signed)


$$\text{if } (![rs]_{31})$$
$$PC_{31..02} \leftarrow \{ PC_{31..02} + (sign-extend-16-to-30 \text{ boffs16}) \}_{29..00}$$
$$PC_{01..00} \leftarrow 0$$

BGEZAL


Branch if greater than or equal to zero and link (signed)



```
if (![rs]31)
    [31] ← { PC + 4 }31..00
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00
```

BGTZ

Branch if greater than zero (signed)




```
if (![rs]31 && ([rs] != 0))
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
```

BL

Branch Long

RawH




```
PC31..02 ← { PC31..02 + (sign-extend-26-to-30 boffs26) }29..00
PC01..00 ← 0
```

BLAL

Branch Long and Link


RawH



```
[31] ← { PC + 4 }31..00
PC31..02 ← { PC31..02 + (sign-extend-26-to-30 boffs26) }29..00
PC01..00 ← 0
```

BLEZ


Branch if less than or equal to zero (signed)



```
if ([rs]31 || ([rs] == 0))
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
```

BLTZ


Branch if less than zero (signed)



```
if ([rs]31)
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
```

BLTZAL


Branch if less than zero and link (signed)



```
if ([rs]31)
    [31] ← { PC + 4 }31..00
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00
```

BNE


Branch if not equal



```
if ([rs] != [rt])
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
```

BNEA

Branch if not equal and add




```

if ([rs] != [rt])
    PC31..02 ← { PC31..02 + (sign-extend-16-to-30 boffs16) }29..00
    PC01..00 ← 0
    [rs] = [rs] + SR[BR_INC]

```

CLZ

Count Leading Zero




```

[rd]05..00 ←  $\sum_{i=0}^{31} ([rs]_{31..i} ? 0 : 1)$ 
[rd]31..06 ← 0

```

DIV

Divide Signed



```


LO ← { [rs] /signed [rt] }31..0 }31..0
HI ← { [rs] %signed [rt] }31..0 }63..32

if ([rt] == 0)
    HI ← [rs]
    if ([rs]31)
        LO ← 1
    else
        LO ← -1

```

DIVU

Divide Unsigned




$$LO \leftarrow \{ [rs] \text{ / } \textit{unsigned} [rt] \}_{31..0} \}_{31..0}$$

$$HI \leftarrow \{ [rs] \% \textit{unsigned} [rt] \}_{31..0} \}_{63..32}$$

```
if ([rt] == 0)
    HI ← [rs]
    LO ← -1
```

J

Jump




$$PC_{31..02} \leftarrow (\textit{zero-extend-16-to-30} \text{ targ16})$$

$$PC_{01..00} \leftarrow 0$$

JAL

Jump and link




$$[31] \leftarrow \{ PC + 4 \}_{31..00}$$

$$PC_{31..02} \leftarrow (\textit{zero-extend-16-to-30} \text{ targ16})$$

$$PC_{01..00} \leftarrow 0$$

JALR

Jump and link through Register




$$[31] \leftarrow \{ PC + 4 \}_{31..00}$$

$$PC_{31..02} \leftarrow [rs]_{31..02}$$

$$PC_{01..00} \leftarrow 0$$

JEQL

Jump if not equal and link




```

if ([rs] == [rt])
    [31] ← { PC + 4 }31..00
    PC31..02 ← (zero-extend-16-to-30 targ16)
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00

```

JGEZL

Jump if greater than or equal to zero and link (signed)




```

if (![rs]31)
    [31] ← { PC + 4 }31..00
    PC31..02 ← (zero-extend-16-to-30 targ16)
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00

```

JGTZL

Jump if greater than zero and link (signed)




```

if (![rs]31 && ([rs] != 0))
    [31] ← { PC + 4 }31..00
    PC31..02 ← (zero-extend-16-to-30 targ16)
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00

```

JLEZL


Jump if less than or equal to zero and link (signed)



```
if ([rs]31 || ([rs] == 0))
    [31] ← { PC + 4 }31..00
    PC31..02 ← (zero-extend-16-to-30 targ16)
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00
```

JLTZL


Jump if less than zero and link (signed)



```
if ([rs]31)
    [31] ← { PC + 4 }31..00
    PC31..02 ← (zero-extend-16-to-30 targ16)
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00
```

JNEL


Jump if not equal and link



```
if ([rs] != [rt])
    [31] ← { PC + 4 }31..00
    PC31..02 ← (zero-extend-16-to-30 targ16)
    PC01..00 ← 0
else
    [31] ← { PC + 4 }31..00
```

JR

Jump through Register




$PC_{31..02} \leftarrow [rs]_{31..02}$
 $PC_{01..00} \leftarrow 0$

JRHOFF

Jump through Register and Disable Hardware ICaching

RawH




$PC_{31..02} \leftarrow [rs]_{31..02}$
 $PC_{01..00} \leftarrow 0$

JRHON

Jump through Register and Enable Hardware ICaching


RawH



$PC_{31..02} \leftarrow [rs]_{31..02}$
 $PC_{01..00} \leftarrow 0$

LB


Load Byte



$ea \leftarrow \{ [base] + (sign-extend-16-to-32\ soffs16) \}_{31..0}$
 $[rt] \leftarrow (sign-extend-8-to-32\ (cache-read-byte\ ea))$

LBU


Load Byte Unsigned



$ea \leftarrow \{ [base] + (sign-extend-16-to-32 \text{ soffs16}) \}_{31..0}$
 $[rt]_{31..8} \leftarrow 0$
 $[rt]_{7..0} \leftarrow (cache-read-byte \text{ } ea)$

LH


Load Halfword



$ea \leftarrow \{ [base] + (sign-extend-16-to-32 \text{ soffs16}) \}_{31..0}$
 $[rt] \leftarrow (sign-extend-16-to-32 \text{ } (cache-read-half-word \text{ } ea))$

LHU

Load Halfword Unsigned



$ea \leftarrow \{ [base] + (sign-extend-16-to-32 \text{ soffs16}) \}_{31..0}$
 $[rt]_{31..16} \leftarrow 0$
 $[rt]_{15..0} \leftarrow (cache-read-half-word \text{ } ea)$

LI

Load Immediate (Assembly Macro)

MIPS


①-②

li rd, uimm32
li! rd, uimm32

$[rd]_{31..0} \leftarrow uimm32_{31..0}$

LW


Load Word



$ea \leftarrow \{ [base] + (sign-extend-16-to-32\ soffs16) \}_{31..0}$
 $[rt] \leftarrow (cache-read-word\ ea)$

MAGIC

User-specified simulator function




$[rt]_{31..00} \leftarrow (user\ function\ code\ [rs])$ - On BTL simulator

$[rt]_{31..00} \leftarrow unspecified\ value$ - On RTL and hardware

MFFD


Move from FD



$[rd]_{31..00} \leftarrow FD_{31..00}$

MFHI


Move from HI



$[rd]_{31..00} \leftarrow HI_{31..00}$

MFLO

Move from LO



$[rd]_{31..00} \leftarrow LO_{31..00}$

MOVE

MOVE (Assembly Macro)


MIPS

①
move rd, rt
move! rd, rt

$[rd]_{31..0} \leftarrow [rt]_{31..0}$

MTFD


Move to FD



$FD_{31..00} \leftarrow [rs]_{31..00}$

MTHI


Move to HI



$HI_{31..00} \leftarrow [rs]_{31..00}$


MTLO

Move to LO




$LO_{31..00} \leftarrow [rs]_{31..00}$

MULLO Multiply Low Signed




$$[rd]_{31..00} \leftarrow \{ [rs]_{31..00} *_{signed} [rt]_{31..00} \}_{31..0}$$

MULLU Multiply Low Unsigned




$$[rd]_{31..00} \leftarrow \{ [rs]_{31..00} *_{unsigned} [rt]_{31..00} \}_{31..0}$$

MULHI Multiply High Signed



$$[rd]_{31..00} \leftarrow \{ [rs]_{31..00} *_{signed} [rt]_{31..00} \}_{63..32}$$


MULHU Multiply High Unsigned



$$[rd]_{31..00} \leftarrow \{ [rs]_{31..00} *_{unsigned} [rt]_{31..00} \}_{63..32}$$

NOR Nor Bitwise

MIPS




$$[rd]_{31..00} \leftarrow \sim([rs]_{31..00} \mid [rt]_{31..00})$$

OR

Or Bitwise


MIPS


$$[rd]_{31..00} \leftarrow [rs]_{31..00} \mid [rt]_{31..00}$$

ORI


Or Bitwise Immediate

MIPS


$$\begin{aligned} [rt]_{31..16} &\leftarrow [rs]_{31..16} \\ [rt]_{15..00} &\leftarrow [rs]_{15..00} \mid \text{uimm16} \end{aligned}$$


POPC

Population Count


$$\begin{aligned} [rd]_{04..00} &\leftarrow \sum_{i=0}^{31} [rs]_i \\ [rd]_{31..05} &\leftarrow 0 \end{aligned}$$


RLM

Rotate Left and Mask


$$\begin{aligned} \text{mask} &\leftarrow (\text{create-mask } mb \text{ me } z) \\ [rt]_{31..0} &\leftarrow (\text{left-rotate } [rs]_{31..0} \text{ ra}) \& \text{mask} \end{aligned}$$

RLMI


Rotate Left and Masked Insert



$\text{mask} \leftarrow (\text{create-mask } \text{mb } \text{me } \text{z})$
 $[\text{rt}]_{31..0} \leftarrow ((\text{left-rotate } [\text{rs}]_{31..0} \text{ ra}) \& \text{mask}) \mid ([\text{rt}]_{31..0} \& \sim \text{mask})$

RLVM

Rotate Left Variable and Mask



$\text{mask} \leftarrow (\text{create-mask } \text{mb } \text{me } \text{z})$
 $\text{ra} \leftarrow [\text{rt}]_{31..0}$
 $[\text{rd}]_{31..0} \leftarrow (\text{left-rotate } [\text{rs}]_{31..0} \text{ ra}) \& \text{mask}$

RRM

Rotate Right and Mask (Assembly Macro)

① *rrm rt, rs, ra, mask*
rrm! rt, rs, ra, mask

$[\text{rt}]_{31..0} \leftarrow (\text{right-rotate } [\text{rs}]_{31..0} \text{ ra}) \& \text{mask}$

(instruction is implemented using RLM; same set of masks are valid)

RRMI

Rotate Right and Mask (Assembly Macro)

① *rrmi rt, rs, ra, mask*
rrmi! rt, rs, ra, mask


$[\text{rt}]_{31..0} \leftarrow ((\text{right-rotate } [\text{rs}]_{31..0} \text{ ra}) \& \text{mask}) \mid ([\text{rt}]_{31..0} \& \sim \text{mask})$

(instruction is implemented using RLMI; same set of masks are valid)

SLL

Shift Left Logical

MIPS




$[rt]_{31..sa} \leftarrow [rs]_{(31-sa)..0}$
 $[rt]_{(sa-1)..0} \leftarrow 0$

SLLV

Shift Left Logical Variable

MIPS




$sa \leftarrow [rt]_{4..0}$
 $[rd]_{31..sa} \leftarrow [rs]_{(31-sa)..0}$
 $[rd]_{(sa-1)..0} \leftarrow 0$

SLT

Set Less Than Signed

MIPS




$[rd]_{00} \leftarrow ([rs]_{31..00} <_{signed} [rt]_{31..00}) ? 1 : 0$
 $[rd]_{31..01} \leftarrow 0$

SLTI

Set Less Than Immediate Signed

MIPS




$simm32 \leftarrow (sign-extend-16-to-32\ simm16)$
 $[rt]_{00} \leftarrow ([rs]_{31..00} <_{signed} simm32_{31..00}) ? 1 : 0$
 $[rt]_{31..01} \leftarrow 0$

SLTIU

Set Less Than Immediate Unsigned

MIPS




$uimm32 \leftarrow (sign-extend-16-to-32 \text{ } simm16)$
 $[rt]_{00} \leftarrow ([rs]_{31..00} <_{unsigned} uimm32_{31..00}) ? 1 : 0$
 $[rt]_{31..01} \leftarrow 0$

SLTU

Set Less Than Unsigned

MIPS




$[rd]_{00} \leftarrow ([rs]_{31..00} <_{unsigned} [rt]_{31..00}) ? 1 : 0$
 $[rd]_{31..01} \leftarrow 0$

SRA

Shift Right Arithmetic

MIPS




$[rt]_{(31-sa)..00} \leftarrow [rs]_{31..sa}$
 $[rt]_{31..(31-sa)} \leftarrow [rs]_{31}$

SRAV

Shift Right Arithmetic Variable

MIPS




$sa \leftarrow [rt]_{4..0}$
 $[rd]_{(31-sa)..00} \leftarrow [rs]_{31..sa}$
 $[rd]_{31..(31-sa)} \leftarrow [rs]_{31}$

SRL

Shift Right Logical


MIPS


$$[rt]_{(31-sa)..0} \leftarrow [rs]_{31..sa}$$
$$[rt]_{(31..(32-sa))} \leftarrow 0$$

SRLV

Shift Right Logical Variable


MIPS


$$sa \leftarrow [rt]_{4..0}$$
$$[rd]_{(31-sa)..00} \leftarrow [rs]_{31..sa}$$
$$[rd]_{31..(32-sa)} \leftarrow 0$$

SUBU

Subtract


MIPS


$$[rd]_{31..00} \leftarrow \{ [rs]_{31..00} - [rt]_{31..00} \}_{31..0}$$

SB

Store Byte

MIPS



$$ea \leftarrow \{ [base] + (sign-extend-16-to-32 \text{ soffs16}) \}_{31..0}$$

(cache-write-byte ea [rt])

SH

Store Halfword

MIPS




$ea \leftarrow \{ [base] + (sign-extend-16-to-32\ soffs16) \}_{31..0}$
(*cache-write-half-word* ea $[rt]$)

SW

Store Word

MIPS




$ea \leftarrow \{ [base] + (sign-extend-16-to-32\ soffs16) \}_{31..0}$
(*cache-write-word* ea $[rt]$)

XOR

Exclusive-Or Bitwise

MIPS




$[rd]_{31..0} \leftarrow [rs]_{31..00} \wedge [rt]_{31..00}$

XORI

Exclusive-Or Bitwise Immediate


MIPS



$[rt]_{31..16} \leftarrow [rs]_{31..16}$
 $[rt]_{15..00} \leftarrow [rs]_{15..00} \wedge uimm16$

ADD.s


Add (single precision floating point)



$$[rd]_{31..0} \leftarrow [rs]_{31..0} +_{IEEE-754} [rt]_{31..0}$$

C.<XXX>.s

Compare (single precision floating point)



$$\begin{aligned} \{invalid_0 \ result_0\} &\leftarrow (floating_point_compare \ \langle XXX \rangle \ [rs]_{31..0} \ [rt]_{31..0}) \\ [rd]_0 &\leftarrow result_0 \\ [rd]_{31..1} &\leftarrow 0 \\ SR[FPSR]_4 &\leftarrow SR[FPSR]_4 \mid invalid_0 \end{aligned}$$

e.g.,


`c.ult.s $4, $5, $7`

The code values of 8..15 correspond to instructions that set the `invalid` bit of the floating point status register (FPSR) when an unordered comparison occurs. The behavior of the helper function *floating-point-compare* matches the MIPS ISA [?] and is shown in the following table:

| Predicate | | | <i>floating-point-compare</i> outputs for each comparison outcome | | | | | |
|-----------|-------------------|-----------------|--|---|----|-----------|----------------------|-----------|
| code | Mnemonic <XXX> | Description | result ₀ | | | | invalid ₀ | |
| | | | > | < | == | unordered | >, <, == | unordered |
| 0 | F | False | ↑ | 0 | 0 | 0 | ↑ | ↑ |
| 1 | UN | Unordered | | 0 | 0 | 1 | | |
| 2 | EQ | Equal | | 0 | 1 | 0 | | |
| 3 | UEQ | Unordered == | | 0 | 1 | 1 | | |
| 4 | OLT | Ordered < | | 1 | 0 | 0 | | 0 |
| 5 | ULT | Unordered or < | | 1 | 0 | 1 | | ↓ |
| 6 | OLE | Ordered ≤ | | 1 | 1 | 0 | | |
| 7 | ULE | Unordered or ≤ | | 1 | 1 | 1 | | ↓ |
| 8 | SF | Signaling False | 0 | 0 | 0 | 0 | 0 | ↑ |
| 9 | NGLE | Not (> or ≤) | | 0 | 0 | 1 | | |
| 10 | SEQ | Signaling == | | 0 | 1 | 0 | | |
| 11 | NGL | Not (< or >) | | 0 | 1 | 1 | | |
| 12 | LT | < | | 1 | 0 | 0 | | |
| 13 | NGE | Not ≥ | | 1 | 0 | 1 | | |
| 14 | LE | ≤ | | 1 | 1 | 0 | | |
| 15 | NGT | Not > | | 1 | 1 | 1 | | ↓ |

CVT.s


Convert from integer to float



$$[rd]_{31..0} \leftarrow (\text{convert-from-integer-to-float } [rs]_{31..0})$$

CVT.w


Convert from float to integer, with round to nearest even



$$[rd]_{31..0} \leftarrow (\text{convert-from-float-to-integer-round-nearest-even } [rs]_{31..0})$$

DIV.s


Divide (single precision floating point)



$$FD_{31..0} \leftarrow [rs]_{31..0} /_{IEEE-754} [rt]_{31..0}$$

MUL.s


Multiply (single precision floating point)



$$[rd]_{31..0} \leftarrow [rs]_{31..0} *_{IEEE-754} [rt]_{31..0}$$

NEG.s


Negate (single precision floating point)



$$[rd]_{31..0} \leftarrow -_{IEEE-754} [rs]_{31..0}$$


SUB.s

Subtract (single precision floating point)


$$[rd]_{31..0} \leftarrow [rs]_{31..0} -_{IEEE-754} [rt]_{31..0}$$


TRUNC.w

Convert from float to integer, with truncation


$$[rd]_{31..0} \leftarrow (\text{convert-from-float-to-integer-truncate } [rs]_{31..0})$$


DRET

Return from User Interrupt


$$\begin{aligned} PC_{31..02} &\leftarrow SR[EX_UPC]_{31..02} \\ PC_{01..00} &\leftarrow 0 \\ SR[EX_BITS]_{31} &\leftarrow 1'b1 \end{aligned}$$


ERET

Return from System Interrupt


$$\begin{aligned} PC_{31..02} &\leftarrow SR[EX_PC]_{31..02} \\ PC_{01..00} &\leftarrow 0 \\ SR[EX_BITS]_{30} &\leftarrow 1'b1 \end{aligned}$$


IHDR

Create Internal Header



ILW


Instruction Load Word



The additional cycle of occupancy is a cycle stolen from the fetch unit on access.


INTOFF

Disable System Interrupts



INTON


Enable System Interrupts



$SR[EX_BITS]_{30} \leftarrow 1'b1$

ISW

Instruction Store Word




$ea_{31..2} \leftarrow \{ [base] + (sign-extend-16-to-32\ soffs16) \}_{31..2}$
 $ea_{1..0} \leftarrow 0$
(proc-imem-store ea [rt])

Steals one fetch cycle from compute processor fetch unit.

MFEC

Move From Event Counter



$[rd]_{31..00} \leftarrow EC[rs]$

Note: MFEC captures its value in the RF stage. This is because the event counters are located physically quite distant from the bypass paths of the processor, so the address is transmitted in RF, and the output given in EXE. For example,

```

lw $0,4($0)           # cache miss in TV stage, pipeline frozen
nop                   # occupies TL stage
mfec $4, EC_CACHE_MISS # EXE stage -- will not register cache miss
mfec $4, EC_CACHE_MISS # RF          -- will register cache miss

```

Additionally, there is one cycle of lag between when the event actually occurs and when the event counter is updated. For example, assuming no outside stalls like cache misses or interrupts,


```

mtec EC_xxx, $4       # write an event counter
mfec $5, EC_xxx       # reads old value
mfec $5, EC_xxx       # reads new value $

```

MFSR

Move From Status / Control Register




$[rd]_{31..00} \leftarrow SR[rs]$

Section 1.4 describes the status registers.

MLK

MDN Lock

RawH




$SR[EX_BITS]_{30} \leftarrow 1'b0;$
(*icache-prefetch* PC uimm5)

Signals to hardware or software caching system that the following uimm5 cache lines needs to be resident in the instruction cache for correct execution to occur. Disables interrupts. This allows instruction sequences to access the memory network without concern that the i-caching system will also access it.

MUNLK

MDN Unlock

RawH




$SR[EX_BITS]_{30} \leftarrow 1'b1$

Marks end of MDN-locked region. Enables interrupts.

MTEC


Move To Event Counter



$EC[rt] \leftarrow [rs]_{31..00}$

MTSR

Move To Status / Control Register




$$SR[rt] \leftarrow [rs]_{31..00}$$

Section 1.4 describes the status registers. Note that not all status register bits are fully writable, so some bits may not be updated as a result of an MTSR instruction.

MTSRi

Move to Status / Control Immediate




$$SR[rt]_{31..16} \leftarrow 0;$$

$$SR[rt]_{15..00} \leftarrow uimm16;$$

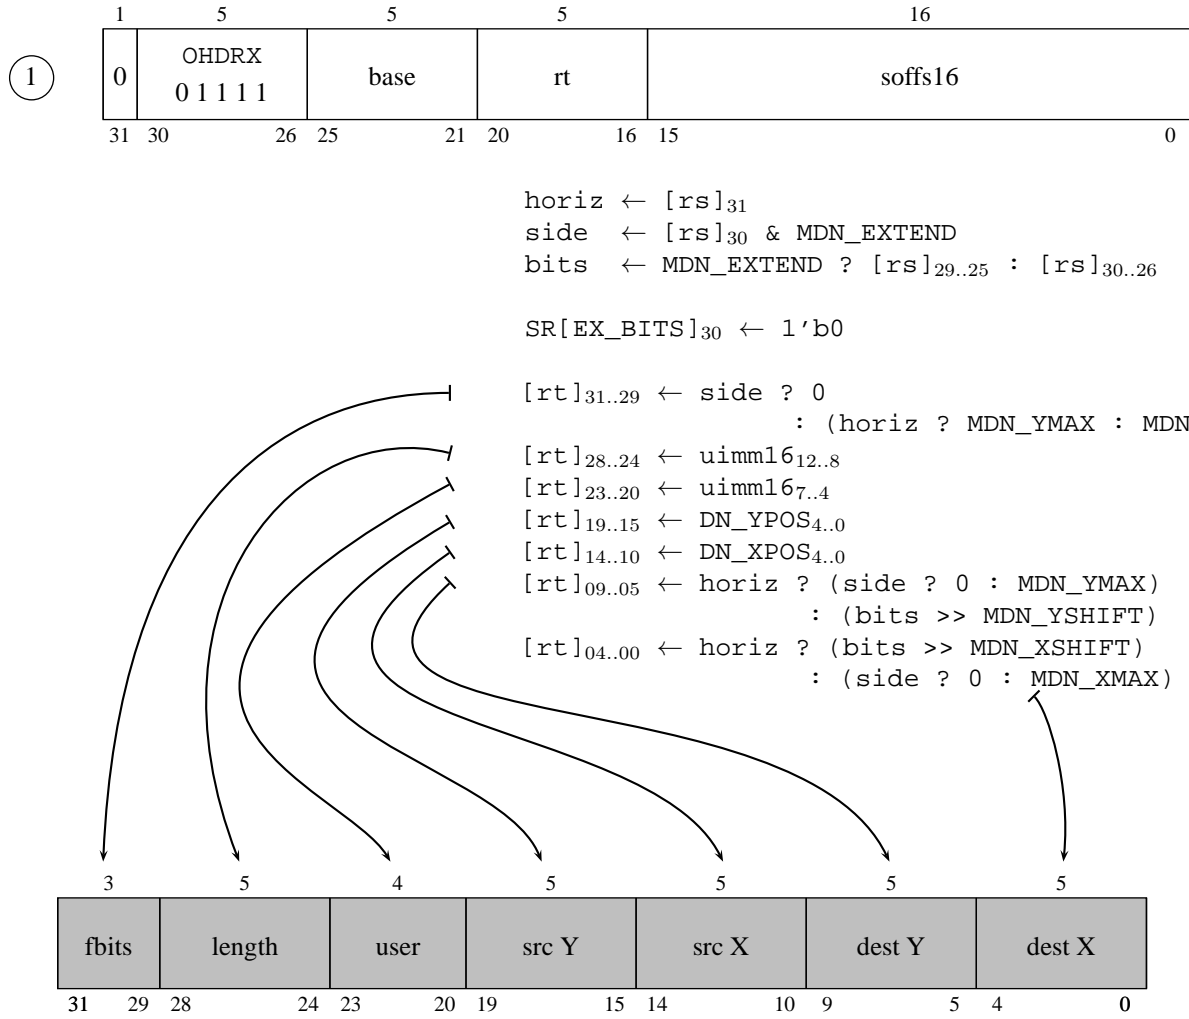
Section 1.4 describes the status registers. Note that not all status register bits are fully writable, so some bits may not be updated as a result of an MTSR instruction.

Create Outside Header



OHDRX


Create Outside Header; Disable System Interrupts



OHDRX takes an address and an immediate field and produces a header suitable for injecting into the MDN. The immediate field specifies the `user` and `length` fields of the message header. OHDRX maps the address to an I/O port, which effectively wraps the address space around the periphery of the chip. Since the MDN must be accessed with interrupts disabled, OHDRX provides a cheap way of doing this. Raw's hardware data cache uses a private copy of this logic to implement Raw's *memory hash function*.

PWRBLK

Power Block




1+

$[rd]_{31} \leftarrow$ data available in cgni
 $[rd]_{30..13} \leftarrow 0;$
 $[rd]_{12} \leftarrow$ data available in cNi
 $[rd]_{11} \leftarrow$ data available in cEi
 $[rd]_{10} \leftarrow$ data available in cSi
 $[rd]_9 \leftarrow$ data available in cWi
 $[rd]_8 \leftarrow$ data available in csti
 $[rd]_7 \leftarrow$ data available in cNi2
 $[rd]_6 \leftarrow$ data available in cEi2
 $[rd]_5 \leftarrow$ data available in cSi2
 $[rd]_4 \leftarrow$ data available in cWi2
 $[rd]_3 \leftarrow$ data available in csti2
 $[rd]_2 \leftarrow$ data available in cmni
 $[rd]_1 \leftarrow$ timer interrupt went off
 $[rd]_0 \leftarrow$ external interrupt went off

Stalls in RF stage until output is non-zero.

SWLW

Switch Load Word



5


1

$ea_{31..2} \leftarrow \{ [base] + (sign-extend-16-to-32 \text{ soffs16}) \}_{31..2}$
 $ea_{1..0} \leftarrow 0$
 $[rt] \leftarrow (static-router-imem-load \text{ } ea)$

Steals one fetch cycle from static router.

SWSW

Switch Store Word



1

$ea_{31..2} \leftarrow \{ [base] + (sign-extend-16-to-32 \text{ soffs16}) \}_{31..2}$
 $ea_{1..0} \leftarrow 0$
 $(static-router-imem-store \text{ } ea \text{ } [rt])$

Steals one fetch cycle from static router.

UINTOFF

Disable User Interrupts

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|----|---------------|----|--|--|-------|----|----|--|-------|--|----|----|-------|--|--|----|-------|--|--|--|-------------------|---|--|--|--|---|---|--|--|--|
| ① | 1 | 5 | | | | 5 | | | | 5 | | | | 5 | | | | 6 | | | | | | | | | | | | | |
| | 0 | COMM 01011 | | | | 00000 | | | | 00000 | | | | 00000 | | | | 00000 | | | | UINTOFF 000010 | | | | | | | | | |
| | 31 | 30 | 26 | | | | 25 | 21 | | | | 20 | 16 | | | | 15 | 11 | | | | 10 | 6 | | | | 5 | 0 | | | |

$SR[EX_BITS]_{31} \leftarrow 1'b0;$

UINTON

Enable User Interrupts

| | | | | | | | | | | | | | |
|---|----|---------------|-------|-------|-------|-------|------------------|----|----|----|---|---|---|
| ① | 1 | 5 | 5 | 5 | 5 | 5 | 6 | | | | | | |
| | 0 | COMM 01011 | 00000 | 00000 | 00000 | 00000 | UINTON 001010 | | | | | | |
| | 31 | 30 | 26 | 25 | 21 | 20 | 16 | 15 | 11 | 10 | 6 | 5 | 0 |

$SR[EX_BITS]_{31} \leftarrow 1'b1;$

1.1.3 Cache Management in Raw and RawH

Because Raw’s memory model is shared memory but not hardware cache-coherent, effective and fast software cache management instructions are essential. One tile may modify a data structure through its caching system, and then want to make it available to a consuming tile or I/O device. To accomplish this in a cache-coherent way, the sender tile must explicitly flush and/or invalidate the data, and then send an MDN Relay message that bounces off the relevant DRAM I/O Port (indicating that all of the memory accesses have reached the DRAM) to the consumer. The consumer then knows that the DRAM has been updated with the correct values.


To provide effective cache management, there are two series of cache management instructions. Both series allow cache lines to be flushed and/or invalidated. The first series, `ainv`, `afl`, and `aflinv`, takes as input a data address. This address, if it is resident in the cache, is translated into a `<set, line>` which is used to identify the physical cache line. The second series of instructions, `tagsw`, `taglv`, `tagla`, and `tagfl`, takes a `<set, line>` pair directly.

The address-based instructions are most effective when the range of addresses residing in the cache is relatively small. If $|A|$ is the size of the address range that needs to be flushed, this series can flush the range in time $\theta(|A|)$.

The tag-based instructions are most effective when the processor needs to invalidate or flush large ranges of address space that exceed the cache size. In this case, the address range can be manipulated faster by using the tag-based instructions to scan the tags of the cache and selectively invalidate and/or flush the contents. In this case, the operations can occur in $\theta(|C|)$, where $|C|$ is the size of the cache. The `tagla` and `taglv` operations allow the cache line tags to be inspected, `tagfl` can be used to flush the contents, and `tagsw` can be used to rewrite (or zero) the tags. Of course, the `tagxxx` series of instructions can accomplish more than simply flushing or invalidating. They provide an easy way to manipulate the cache state directly for verification purposes and boot-time cache initialization.

AINV

Address Invalidate




$ea \leftarrow [rs] + (S \ll 14) + (AAAA \ll kDataCacheLineSize)$

if (*cache-contains* ea)

 TAGS[(*cache-get-tag* ea)].valid $\leftarrow 0$ # stall 4 cycles

AFL

Address Flush




```
ea ← [rs] + (S << 14) + (AAAA << kDataCacheLineSize)
```

```
if (cache-contains ea)
{
  <set,line> ← (cache-get-tag ea)
  TAGS[<set,line>].mru ← !set

  if (TAGS[<set,line>].dirty)
  {
    TAGS[<set,line>].dirty ← 0
    (cache-copy-back <set,line>) # stall >= 13 cycles
  }
  else
    # stall 5 cycles
}
```

AFLINV

Address Flush and Invalidate




```
ea ← [rs] + (S << 14) + (AAAA << kDataCacheLineSize)
```

```
if (cache-contains ea)
{
  <set,line> ← (cache-get-tag ea)
  if (TAGS[<set,line>].dirty)
  {
    TAGS[<set,line>].dirty ← 0
    TAGS[<set,line>].valid ← 0
    (cache-copy-back ea) # stall >= 13 cycles
  }
  else
    # stall 5 cycles
}
```

TAGFL

Tag Flush



```


set      ← [rs]14 ^ S
line8..0 ← [rs]13..5 + AAAA

if (TAGS[<set,line>].valid)
{
  TAGS[<set,line>].mru ← !set
  if (TAGS[<set,line>].dirty)
  {
    TAGS[<set,line>].dirty ← 0
    (cache-copy-back <set,line>)    # stall >= 13 cycles
  }
  else
    # stall 5 cycles
}

```

TAGLA

Tag Load Address



```


set      ← [rs]14 ^ S
line8..0 ← [rs]13..5 + AAAA

[rd] ← { TAGS[<set,line>].addr17..00 line8..0 [rs]4..0 }

```

TAGLV

Tag Load Valid



```


set      ← [rs]14 ^ S
line8..0 ← [rs]13..5 + AAAA

[rd] ← TAGS[<set,line>].valid

```

TAGSW

Tag Store Word



$set \leftarrow [rs]_{14} \wedge S$
 $line_{8..0} \leftarrow [rs]_{13..5} + AAAA$

$TAGS[<set,line>].valid_0 \leftarrow [rt]_{18}$
 $TAGS[<set,line>].addr_{17..00} \leftarrow [rt]_{17..00}$
 $TAGS[<set,line>].dirty_0 \leftarrow 0$

Should not be issued the cycle after a load or store instruction because of write-after-write hazards on the tag memory.

1.2 Semantic Helper Functions

This section gives the semantics of the helper-functions used in the previous section. This thesis uses little-endian bit-ordering exclusively.

| | | |
|---|---|---|
| $w_{x..y}$ | → | Bits $x..y$, inclusive, of w . If $(x < y)$, the empty string. |
| $\{ w \ z \}$ | → | Concatenate the bits of w and z together. w will occupy the more significant bits. |
| z^n | → | Concatenate n copies of z together. |
| $(\text{sign-extend-16-to-32} \ \text{simm16})$ | → | $\{ (\text{simm16}_{15})^{16} \ \text{simm16}_{15..00} \}$ |
| $(\text{sign-extend-26-to-30} \ \text{simm26})$ | → | $\{ (\text{simm26}_{25})^4 \ \text{simm26}_{25..00} \}$ |
| $(\text{sign-extend-16-to-30} \ \text{simm16})$ | → | $\{ (\text{simm16}_{15})^{14} \ \text{simm16}_{15..00} \}$ |
| $(\text{zero-extend-16-to-32} \ \text{uimm16})$ | → | $\{ 0^{15..0} \ \text{uimm16}_{15..00} \}$ |
| $(\text{left-rotate} \ \text{uimm32} \ ra)$ | → | $\{ \text{uimm32}_{(31-ra)..0} \ \text{uimm32}_{31..(32-ra)} \}$ |
| $(\text{right-rotate} \ \text{uimm32} \ ra)$ | → | $\{ \text{uimm32}_{(ra-1)..0} \ \text{uimm32}_{31..ra} \}$ |
| $(\text{cache-contains} \ \text{addr})$ | → | Returns 1 if valid cache line corresponding to addr is in cache, otherwise 0. |
| $(\text{cache-get-tag} \ \text{addr})$ | → | Returns $\langle \text{set}, \text{line} \rangle$ pair corresponding to addr in cache. |
| $(\text{cache-copy-back} \ \text{tagid})$ | → | Sends update message containing data corresponding to tagid to owner DRAM. |
| $(\text{cache-read-byte} \ \text{addr})$ | → | Ensure cache line corresponding to addr is in cache, return byte at addr . |
| $(\text{cache-read-half-word} \ \text{addr})$ | → | Ensure cache line corresponding to addr is in cache, return half-word at $\{ \text{addr}_{31..1} \ 0^1 \}$. |
| $(\text{cache-read-word} \ \text{addr})$ | → | Ensure cache line corresponding to addr is in cache, return word at $\{ \text{addr}_{31..2} \ 0^2 \}$. |
| $(\text{cache-write-byte} \ \text{addr} \ \text{val})$ | → | Ensure cache line corresponding to addr is in cache, write $\text{val}_{7..0}$ to addr . |
| $(\text{cache-write-half-word} \ \text{addr} \ \text{val})$ | → | Ensure cache line corresponding to addr is in cache, write $\text{val}_{15..0}$ to $\{ \text{addr}_{31..1} \ 0^1 \}$. |
| $(\text{cache-write-word} \ \text{addr} \ \text{val})$ | → | Ensure cache line corresponding to addr is in cache, write $\text{val}_{31..0}$ to $\{ \text{addr}_{31..2} \ 0^2 \}$. |
| $(\text{create-mask} \ mb \ me \ z)$ | → | if (z) if ($me_{1..0} == 0b00$) $\{ mb_{4..0} \ me_{4..2} \}^4$ if ($me_{1..0} == 0b11$) $\{ mb_4^4 \ mb_3^4 \ mb_2^4 \ mb_1^4 \ mb_0^4 \ me_4^4 \ me_3^4 \ me_2^4 \}$ else if ($mb \leq_{\text{unsigned}} me$) $\{ 0^{31..(me+1)} \ 1^{me..mb} \ 0^{(mb-1)..0} \}$ else |

$$\{ 1^{31..(mb+1)} 0^{mb..me} 1^{(me-1)..0} \}$$

The last line was a specification bug as it does not generate every mask with a single zero. A better version is:

$$\{ 1^{31..(mb+1)} 0^{mb..me+1} 1^{(me)..0} \}$$

(icache-prefetch addr lines) → Ensure *lines* instruction cache lines following cache line containing *addr* are resident in instruction cache.

(static-router-imem-store addr data) → Writes 32-bit value *data* into static router instruction cache at location *addr*.

(static-router-imem-load addr data) → Loads 32-bit value *data* from static router instruction cache at location *addr*.

(proc-imem-store addr data) → Writes 32-bit value *data* into static router instruction cache at location *addr*.

(proc-imem-load addr data) → Loads 32-bit value *data* from static router instruction cache at location *addr*.

1.3 Opcode Maps

Below are opcode maps which document the allocation of instruction encoding space.

1.3.1 High-Level (“Opcode”) Map

(Instructions with bits 31..29 set to 1 are predicted taken.) Bang instructions all have bit 31 set. however RLM, RLMI and RLVM use bit 28 to indicate bang.

| bits 31..29 | bits 28..26 | | | | | | | |
|----------------|-----------------|-------|------|------------|-------|--------|-------------|--------|
| | 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
| 111 | REGIMM+ | BNEA+ | BNE+ | BEQ+ | BL | BLAL | JNEL+ | JEQL+ |
| 110 | LB! | LBU! | LH! | LHU! | LW! | SLTIU! | SLTI! | ADDIU! |
| 101 | RLM | RLMI | RLVM | | RLM! | RLMI! | RLVM! | |
| 100 | SPECIAL! | ILW! | ORI! | XORI! | ANDI! | SWLW! | FPU! | AUI! |
| 011 | REGIMM- | BNEA- | BNE- | BEQ- | MTSRI | IHDR | JNEL- | JEQL- |
| 010 | LB | LBU | LH | LHU | LW | SLTIU | SLTI | ADDIU |
| 001 | SB | ISW | SH | COM | SW | SWSW | OHDR | OHDRX |
| 000 | SPECIAL | ILW | ORI | XORI | ANDI | SWLW | FPU | AUI |

1.3.2 SPECIAL Submap

(Applies when bits 31..26 are SPECIAL or SPECIAL!)

| bits 5..3 | bits 2..0 | | | | | | | |
|--------------|-----------|-------|-------|--------|------|------|------|------|
| | 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
| 000 | SLL | MAGIC | SRL | SRA | SLLV | | SRLV | SRAV |
| 001 | JR | JALR | JRHON | JRHOFF | | | | |
| 010 | MFHI | MTHI | MFLO | MTLO | MFFD | MTFD | | |
| 011 | MULLO | MULLU | DIV | DIVU | | | | |
| 100 | | ADDU | | SUBU | AND | OR | XOR | NOR |
| 101 | MULHI | MULHU | SLT | SLTU | | | | |
| 110 | | | | | | | | |
| 111 | POPC | CLZ | | | | | | |

1.3.3 FPU Submap

(Applies when bits 31..26 are FPU or FPU!)

| bits | bits 2..0 | | | | | | | |
|------|-----------|--------|-------|-------|-------|---------|-------|-------|
| 5..3 | 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
| 000 | ADD.s | SUB.s | MUL.s | DIV.s | | ABS.s | | NEG.s |
| 001 | | | | | | TRUNC.s | | |
| 010 | | | | | | | | |
| 011 | | | | | | | | |
| 100 | CVT.s | | | | CVT.w | | | |
| 101 | | | | | | | | |
| 110 | C.F | C.UN | C.EQ | C.UEQ | C.OLT | C.ULT | C.OLE | C.ULE |
| 111 | C.SF | C.NGLE | C.SEQ | C.NGL | C.LT | C.NGE | C.LE | C.NGT |

1.3.4 COM Submap

(Applies when bits 31..26 are COM)

| bits | bits 2..0 | | | | | | | |
|------|-----------|--------|---------|-------|-----|--------|------|-----|
| 5..3 | 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
| 000 | DRET | INTOFF | UINTOFF | ERET | | | | |
| 001 | | INTON | UINTON | | | | | |
| 010 | MFSR | MTSR | MFEC | MTEC | | | | |
| 011 | TAGSW | TAGFL | TAGLA | TAGLV | AFL | AFLINV | AINV | |
| 100 | PWRBLK | | | | | | | |
| 101 | | | | | | | | |
| 110 | | | | | | | | |
| 111 | | | | | | | | |

1.3.5 REGIMM Submap

(Applies when bits 31..26 are REGIMM+ or REGIMM-.) Bit 20 indicates a link instruction, and bit 18 indicates an absolute jump. The conditions are mirrored across these axes when appropriate.

| bits | bits 18..16 | | | | | | | |
|--------|-------------|------|--------|------|-------|-------|-------|-------|
| 20..19 | 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
| 00 | BLTZ | BLEZ | BGEZ | BGTZ | | | | |
| 01 | | | | | J | | | |
| 10 | BLTZAL | | BGEZAL | | JLTZL | JLEZL | JGEZL | JGTZL |
| 11 | | | | | JAL | | | |

1.4 Status and Control Registers

| # | Status Reg Name | R/W | Purpose |
|---|-----------------|-----|---|
| 0 | SW_FREEZE | RW | Switch Processor is Frozen. [00] ← (1 Frozen) (0 Running) |
| 1 | SW_BUF1 | R | # of elements in static router crossbar 1 NIBs [22:20] number of elements in c21 (≤ 4) [19:17] number of elements in cNi (≤ 4) [16:14] number of elements in cEi (≤ 4) [13:11] number of elements in cSi (≤ 4) [10:08] number of elements in cWi (≤ 4) [07:05] number of elements in csti (≤ 4) [04:00] number of elements in csto (≤ 8) |
| 2 | SW_BUF2 | R | # of elements in static router crossbar 2 NIBs [22:20] number of elements in c12 (≤ 4) [19:17] number of elements in cNi ₂ (≤ 4) [16:14] number of elements in cEi ₂ (≤ 4) [13:11] number of elements in cSi ₂ (≤ 4) [10:08] number of elements in cWi ₂ (≤ 4) [07:05] number of elements in csti ₂ (≤ 4) [04:00] number of elements in csto (≤ 8) |
| 3 | MDN_BUF | R | # of elements in MDN router NIBs [19:17] number of elements in cNi (≤ 4) [16:14] number of elements in cEi (≤ 4) [13:11] number of elements in cSi (≤ 4) [10:08] number of elements in cWi (≤ 4) [07:05] number of elements in cmni (≤ 4) [04:00] number of elements in cmno (≤ 16) |
| 4 | SW_PC | RW | Current PC of switch processor. Byte address aligned to eight-byte boundaries. Used primarily for context switching. Generally, writing to this register is used for context-switching purposes. It should only be performed when the switch is FROZEN or if the compute processor program knows absolutely that the switch is stalled at a known PC. Otherwise, the program can no longer assume the static ordering of operands on the SON. Writing to this register causes a branch misprediction in the switch. Allow at least three cycles for corresponding instruction to be executed. |
| 5 | BR_INCR | RW | Signed 32-bit increment value for BNEA instruction. Caller-saved. |

| # | Status Reg Name | | Purpose |
|----|-----------------|----|--|
| 6 | EC_DYN_CFG | RW | <p>Configuration for Event Counting of Dynamic Network events</p> <p>[30:28] Memory Network North (D=N) Configuration</p> <p>[27:25] Memory Network East (D=E) Configuration</p> <p>[24:22] Memory Network South (D=S) Configuration</p> <p>[21:19] Memory Network West (D=W) Configuration</p> <p>[18:16] Memory Network Proc (D=P) Configuration</p> <p>[14:12] Memory Network North (D=N) Configuration</p> <p>[11:09] Memory Network East (D=E) Configuration</p> <p>[08:06] Memory Network South (D=S) Configuration</p> <p>[05:03] Memory Network West (D=W) Configuration</p> <p>[02:00] Memory Network Proc (D=P) Configuration</p> <p>Settings:</p> <p>0 # of cycles output port D wants to transmit but could not because neighbor tile's input buffer is full.</p> <p>1 # of words transmitted from input port D to output port P</p> <p>2 # of words transmitted from input port D to output port W</p> <p>3 # of words transmitted from input port D to output port S</p> <p>4 # of words transmitted from input port D to output port E</p> <p>5 # of words transmitted from input port D to output port N</p> <p>6 # of words transmitted from input port D</p> <p>7 # cycles input port D had data to transmit but was not able to</p> |
| 7 | WATCH_VAL | RW | [31:00] 32-bit timer; increments each cycle |
| 8 | WATCH_MAX | RW | [31:00] value to fire timer interrupt and then zero WATCH_VAL |
| 9 | WATCH_SET | RW | <p>[00] zero WATCH_VAL if <code>cgno</code> is empty or a value was dequeued</p> <p>[01] zero WATCH_VAL if processor issues an instruction</p> |
| 10 | CYCLE_HI | RW | [31:00] high 32-bits of cycle counter |
| 11 | CYCLE_LO | RW | <p>[31:00] low 32-bits of cycle counter</p> <p>Note: To read the cycle counter efficiently, read CYCLE_HI, then CYCLE_LO, then subtract one from CYCLE_LO. Cycle counters are writable to make tests reproducible.</p> |
| 12 | EVENT_CFG2 | RW | [24:0] configures the set of events that causes <code>c_trigger</code> event counters to be incremented. See 1.5. |
| 13 | GDN_RF_VAL | RW | <p>[31:00] GDN refill value</p> <p>When <code>EX_MASK[GDN_REFILL]</code> is enabled, a read from <code>\$cgno</code> will return <code>GDN_RF_VAL</code>, signal an interrupt by setting <code>EX_BITS[GDN_REFILL]</code>, and leave <code>cgno</code> unchanged. This allows <code>cgno</code> to be virtualized, e.g. for context switches and deadlock recovery.</p> |

| # | Status Reg Name | | Purpose |
|----|-----------------|----|---|
| 14 | GDN.REMAIN | RW | [04:00] Number of words remaining to be sent to complete current message on cgno . GDN.COMPLETE interrupt fires when value transitions to zero. OS typically initializes this with GDN.PENDING value to allow GDN messages to complete when context switching. |
| 15 | EX.BASE.ADDR | RW | [31:00] Pointer to beginning of exception vector table. Set to zero at boot time. Applies to RawH . |
| 16 | GDN.BUF | R | <p># of elements in GDN router NIBs</p> <p>[24:20] GDN.PENDING number of elements (≤ 31) that need to be sent to cgno from processor pipeline to complete current message.</p> <p>Note this count does not include those instructions currently in the pipeline; the operating system should flush the pipeline before reading this value. The OS loads this value into the GDN.REMAIN SPR for the GDN.PENDING interrupt to trigger on.</p> <p>[19:17] number of elements (≤ 4) in cNi [16:14] number of elements (≤ 4) in cEi [13:11] number of elements (≤ 4) in cSi [10:08] number of elements (≤ 4) in cWi [07:05] number of elements (≤ 4) in cgni [04:00] number of elements (≤ 16) in cgno</p> |
| 17 | GDN.CFG | RW | <p>General Dynamic Network Configuration</p> <p>[31:27] GDN.XMASK - Masks X bits from an address [26:22] GDN.YMASK - Masks Y bits from an address [21:17] GDN.XADJ - Adjusts from local to global X address [16:12] GDN.YADJ - Adjusts from local to global Y address [11:09] GDN.YSHIFT - Gets Y bits from an address</p> <p>See IHDR instruction.</p> |

| # | Status Reg Name | | Purpose |
|----|-----------------|----|--|
| 18 | STORE_METER | RW | <p>STORE_ACK counters</p> <p>[31:27] PARTNER_Y - Y location of partner port [26:22] PARTNER_X - X location of partner port [21] ENABLE - enable store meter-based stalls [10] DECREMENT_MODE (see below; reads always zero) [9:5] COUNT_PARTNER - # of partner accesses left [4:0] COUNT_NON_PARTNER - # of non-partner accesses left</p> <p>Since the counts are updated as STORE_ACK messages are received over the MDN, care must be taken to update STORE_METER in a way that avoids race conditions.</p> <p>Ordinarily, the only way to do this is to modify the register only when all store-acks have been received.</p> <p>Alternatively, the user may write to the register with DECREMENT_MODE set; in this case the COUNT_NON_PARTNER will be decremented if bit 0 is set, and COUNT_PARTNER will be decremented if bit 5 is set. No other bits are changed. This handles the case where the user is directly transmitting memory packets over the MDN using explicit accesses to <code>cmno</code>, and needs to update the the STORE_ACK counters to reflect this.</p> |
| 19 | MDN_CFG | RW | <p>Memory Dynamic Network Configuration</p> <p>[31:27] DN_XPOS - Absolute X position of tile in array [26:22] DN_YPOS - Absolute Y position of tile in array [21:17] MDN_XMAX - X Coord of East-Most Tiles [16:12] MDN_YMAX - Y Coord of South-Most Tiles [11:09] MDN_XSHIFT - Shift Amount X [08:06] MDN_YSHIFT - Shift Amount Y [00:00] MDN_EXTEND - Use all four edge of chip.</p> <p>These SPRs are used to determine Raw's <i>memory hash function</i> as described in MBT's PhD thesis. This function determines where the data caches send their messages for cache fills and evictions. It also determines the functionality of the OHDR and OHDRX instructions.</p> |
| 20 | EX_PC | RW | PC where system-level exception occurred. |
| 21 | EX_UPC | RW | PC where user-level exception occurred. (GDN_AVAIL is the only user-level exception) |

| # | Status Reg Name | | Purpose |
|----|-----------------|----|---|
| 22 | FPSR | RW | <p>Floating Point Status Register</p> <p>[5] Unimplemented [4] Invalid [3] Divide by Zero [2] Overflow [1] Underflow [0] Inexact operation</p> <p>These bits are sticky; i.e. floating point operations can set but cannot clear these bits. However, the user can freely change the bits via MTSR or MFSR.</p> <p>These flags are set the cycle after the floating point instruction finishes execution; i.e., you need three nops inbetween the last floating point operation and a MFSR to read the correct value.</p> |
| 23 | EVENT_BITS | R | [15:0] the list of events that have triggered |
| 24 | EX_BITS | R | <p>Interrupt Status</p> <p>[31] USER - all user interrupts masked if 0 [30] SYSTEM - all interrupts masked if 0</p> <p>The above can be set/cleared using inton, intoff, uinton, uintoff.</p> <p>[6] EVENT_COUNTER [5] GDN_AVAIL [4] TIMER [3] EXTERNAL [2] TRACE [1] GDN_COMPLETE [0] GDN_REFILL</p> <p>For bits 0..6, a “1” indicates a request for a given interrupt occurred but that it has not yet been serviced.</p> |
| 25 | EX_MASK | RW | <p>Interrupt Mask</p> <p>[6] EVENT_COUNTER [5] GDN_AVAIL [4] TIMER [3] EXTERNAL [2] TRACE [1] GDN_COMPLETE [0] GDN_REFILL</p> <p>A “0” indicates that the exception is suppressed.</p> |

| # | Status Reg Name | | Purpose |
|----|-----------------|----|---|
| 26 | EVENT_CFG | RW | <p>Event Counter Configuration</p> <p>[31:16] Enables for events 16..0 [15:01] PC to profile (omit low two bits) for single mode [00] ← (1 Single Instruction Mode) (0 Global Instruction Mode)</p> |
| 27 | POWER_CFG | RW | <p>Power Saving Configuration</p> <p>[00] Disable comparator toggle-suppression [01] Disable ALU toggle-suppression [02] Disable FPU toggle-suppression [03] Disable Multiplier toggle-suppression [04] Disable Divider toggle-suppression [05] Disable Data Cache toggle-suppression [06] Enable Instruction Memory power saving [07] Enable Data Memory power saving [08] Enable Static Router Memory power saving [09] Disable pwrblk wake up after TIMER interrupt [10] Disable pwrblk wake up after EXTERNAL interrupt [11] Timer wakeup pending on return to pwrblk [12] External wakeup pending on return to pwrblk</p> <p>At reset, POWER_CFG is set to zero. Bits 11 and 12 are set by the processor if the corresponding interrupt is taken while waiting on a pwrblk.</p> |
| 28 | TN_CFG | W | Test Network Configuration |
| 29 | TN_DONE | W | Signal “DONE” on Test Network with value [31:0] |
| 30 | TN_PASS | W | Signal “PASS” on Test Network with value [31:0] |
| 31 | TN_FAIL | W | Signal “FAIL” on Test Network with value [31:0] |

1.5 Event Counting Support

The event counters provide a facility to monitor, profile, and respond to events on a Raw tile. Each tile has a bank of 16 `c.trigger` modules. Each `c.trigger` has a 32-bit counter. These counters count down every time a particular event occurs. The `EVENT_CFG2` register is used to determine which events each `c.trigger` responds to. When the counter transitions from 0 to -1, it will assert a line (the “trigger”) which will hold steady until the user writes a new value into the counter. These triggers are visible in the `EVENT_BITS` register, and are OR’d together to form the `EX_BITS EVENT_COUNTER` bit, which can cause an interrupt. When the trigger is asserted, the `c.trigger` module latches the PC (without the low zero bits) of the instruction that caused the event into bits [31:16] of the counter (the `r1m` instruction can be used to extract them efficiently). The `c.trigger` module will continue to count down regardless of the setting of the trigger. Because the PC is stored in the high bits, there is a window of time in which subsequent events will not corrupt the captured PC. Note that if the event is not instruction related, the setting of the PC in the `c.trigger` is undefined. The event counters can be both read and written by the user. There is typically a one cycle delay between when an event occurs and when an `mfec` instruction will observe it; there is also a delay of two cycles before an event trigger interrupt will fire.

| c.trigger # EVENT_CFG2 Stage | | | Function | Notes |
|------------------------------------|----------|---|--------------------|---|
| 0 | [25] ← 0 | @ | Cycle Count | So handler can bound sampling window. For poor man’s shared memory support. Detects when a resident cache line is marked dirty by a <code>sw</code> to an odd address for the first time. Note: If the <code>sw</code> is preceded by a <code>lw</code> / <code>sw</code> / <code>flush</code> this mechanism does not have the bandwidth to verify the previous state of the bits. It will conservatively count it as an event. |
| 0 | [25] ← 1 | F | Write Over Read | |
| 1 | | M | Cache Writebacks | Includes flushes. |
| 2 | | M | Cache Fills | |
| 3 | | M | Cache Stall Cycles | Total number of cycles that the backend of the pipeline is frozen by the cache state machine. Includes write-back and fill time, as well as time stolen by non-dirty <code>flush</code> instructions. |
| 4 | [0] ← 0 | E | Cache Miss Ops | Number of <code>flush</code> , <code>lw</code> , <code>sw</code> instructions issued. |
| 4 | [0] ← 1 | E | FPU Ops | Number of FPU instructions issued. Includes <code>.s</code> and <code>.w</code> instructions. |

| c.trigger # EVENT_CFG2 Stage | | | Function | Notes |
|------------------------------------|---------|---|-----------------------------|--|
| 5 | [1] ← 0 | E | Possible Mispredicts | Conditional Jumps and Branches, ERET, DRET, JR, JALR. |
| 5 | [1] ← 1 | E | Possible Mispredicts | Possible mispredicts due to wrong SBIT (i.e., only conditional jumps and branches) |
| 6 | [2] ← 0 | E | Actual Mispredicts | Branch mispredictions. |
| 6 | [2] ← 1 | E | Actual Mispredicts | Mispredictions due to wrong SBIT |
| 7 | | @ | Switch Stalls | On static router (Trigger captures static router PC) |
| 8 | | @ | Possible Mispredicts | On static router (Trigger captures static router PC) |
| 9 | | @ | Actual Mispredicts | On static router (Trigger captures static router PC) |
| 10 | | @ | Pseudo Random LFSR | $X_{next} = (X \gg 1) \mid (\text{xor}(X[31,30,10,0]) \ll 31)$ Note: Sampling this more than once per 32 cycles produces highly correlated numbers. |
| 11 | [3] | R | Functional Unit Stalls | Stalls due to bypassing (e.g., the output of a preceding instruction is not available yet) or because of interlocks on the fp/int dividers. |
| 11 | [4] | @ | GP | GDN Processor Port Counting |
| 11 | [5] | @ | MP | MDN Processor Port Counting |
| 11 | [23] | @ | Instructions Issued | # of instructions that enter Execute stage. |
| 12 | [6] | R | Non-cache stalls | # of stalls not due to cache misses. Includes <code>ilw/isw</code> ; if trigger fires on <code>isw/ilw</code> PC will be the PC of the instruction in the RF stage, rather than the <code>ilw/isw</code> instruction. |
| 12 | [7] | @ | GW | GDN West Port Counting |
| 12 | [8] | @ | MW | MDN West Port Counting |
| 13 | [9] | R | <code>ilw/isw</code> | # of <code>ilw/isw</code> instructions issued. |
| 13 | [10] | @ | GS | GDN South Port Counting |
| 13 | [11] | @ | MS | MDN South Port Counting |
| 13 | [24] | @ | Instructions Issued | # of instructions that enter Execute stage. |
| 14 | [12] | R | <code>\$csto</code> stalls | Instruction issue blocked on <code>\$csto</code> full |
| 14 | [13] | R | <code>\$cgno</code> stalls | Instruction issue blocked on <code>\$cgno</code> full |
| 14 | [14] | R | <code>\$cmno</code> stalls | Instruction issue blocked on <code>\$cmno</code> full |
| 14 | [15] | @ | GE | GDN East Port Counting |
| 14 | [16] | @ | ME | MDN East Port Counting |
| 15 | [17] | R | <code>\$csti</code> stalls | Instruction issue blocked on <code>\$csti</code> empty |
| 15 | [18] | R | <code>\$csti2</code> stalls | Instruction issue blocked on <code>\$csti2</code> empty |
| 15 | [19] | R | <code>\$cgni</code> stalls | Instruction issue blocked on <code>\$cgni</code> empty |
| 15 | [20] | R | <code>\$cmni</code> stalls | Instruction issue blocked on <code>\$cmni</code> empty |
| 15 | [15] | @ | GN | GDN North Port Counting |
| 15 | [16] | @ | MN | MDN North Port Counting |

The previous table describes the events that the `c.trigger` modules can be configured to count. The `EVENT_CFG2` column specifies the bit number of `EVENT_CFG2` that must be set in order to enable counting of that event.

The low bits of `EVENT_CFG` allow the user to count events that occurs on a particular instruction at a particular PC instead of across all PCs. For this “single instruction mode”, `EVENT_CFG[0]` is set to 1, and the PC to sample is placed into `EVENT_CFG[15:1]`. In cases where the event does not have an associated main processor PC (marked with the “@” in the table), the `EVENT_CFG` single instruction mode setting is ignored. The high bits of `EVENT_CFG` selectively enable counting on a per event basis, but do not suppress existing triggers.

The `EVENT_CFG2` SPR allows the user to configure the events that a particular `c.trigger` module counts. In some cases multiple enabled events may be connected to the same trigger. In that case, the counters increments each cycle if any such enabled events has occurred. In some cases, there are nonsensical combinations that can be enabled (say `GE` and `$csto` stalls).

The meaning of the `GN`, `GE`, `GS`, `GW`, `GP`, `MN`, `ME`, `MS`, `MW`, and `MP` events are configured by the `EC_DYN_CFG` status/control register. Each event corresponds to a network `N` (`G` = general, `M` = memory) and a direction `D` (`N`=north, `E`=east, ...). The encodings are shown in the table in Section 1.4.

1.6 Exception Vectors

| # | Name | Offset | Purpose |
|---|--------------------|--------|--------------------------|
| 0 | VEC_GDN_REFILL | 0x00 | Dynamic Refill Exception |
| 1 | VEC_GDN_COMPLETE | 0x10 | GDN Send Is Complete |
| 2 | VEC_TRACE | 0x20 | Trace Interrupt |
| 3 | VEC_EXTERN | 0x30 | External Interrupt (MDN) |
| 4 | VEC_TIMER | 0x40 | Timer Exception |
| 5 | VEC_GDN_AVAIL | 0x50 | Data Avail on GDN |
| 6 | VEC_EVENT_COUNTERS | 0x60 | Event Counter Interrupt |

In the Raw architecture, the exception vectors are stored starting at offset zero in instruction memory. In RawH, the exception vectors are stored relative to `SR[EX_BASE_ADDR]`. When an exception occurs, the processor starts fetching from the corresponding exception location. Thus, a `TIMER` exception would start fetching at address `SR[EX_BASE_ADDR] + 0x40`.

Each exception has 4 contiguous instructions; this is enough to do a small amount of work; such as save a register, load a jump address, and branch there:

```
sw $3, interrupt_save($gp)
lw $3, gdn_vec($gp)
jr $3
```